

Data Mining in Python

Witek ten Hove

10/24/22

Table of contents

Preface	3
Prerequisites	3
Purpose of this course	4
Structure of the course	4
About the author	5
1 Setting up your data science environment	6
1.1 Working with Quarto	6
1.2 Working with Git and Github	6
1.3 Using Python virtual environments	6
2 Lazy learning with k-Nearest Neighbors	7
2.1 Business Case: Diagnosing Breast Cancer	7
2.2 Data Understanding	7
2.3 Preparation	8
2.4 Modeling and Evaluation	12
3 Probabilistic Learning with Naive Bayes Classification	15
3.1 Business Case: Filtering Spam	15
3.2 Data Understanding	15
3.3 Preparation	18
3.4 Modeling and Evaluation	18
References	21

Preface

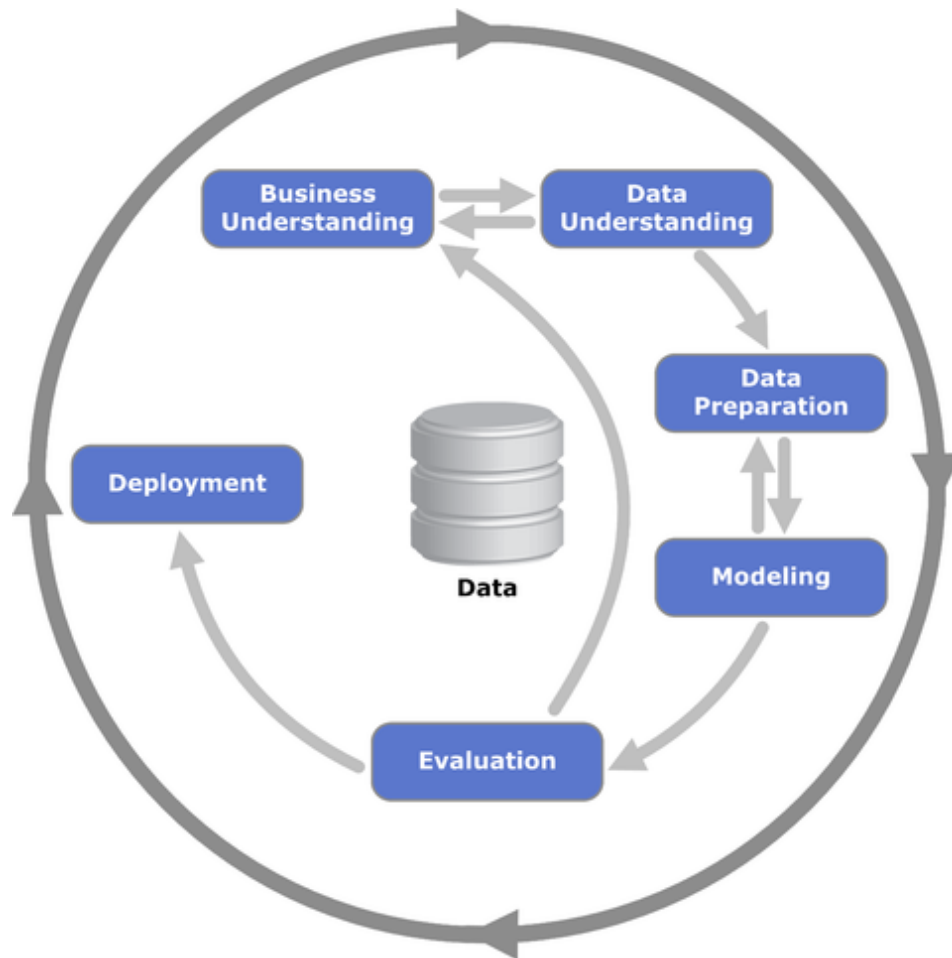


Figure 1: CRISP-DM Model taken from: https://commons.wikimedia.org/wiki/File:CRISP-DM_Process_Diagram.png

Prerequisites

Before starting this module make sure you have:

- access to the book *Provost, F., & Fawcett, T. (2013). Data Science for Business: What you need to know about data mining and data-analytic thinking. O'Reilly Media, Inc.*
- installed [Anaconda](#)
- a [Github](#) account

Purpose of this course

The general learning outcome of this course is:

The student is able to perform a well-defined task independently in a relatively clearly arranged situation, or is able to perform in a complex and unpredictable situation under supervision.

The course will provide you with a few essential data mining skills. The focus will lie on non-linear modeling techniques - k-Nearest Neighbors (kNN) and Naive Bayes classification.

After a successful completion of the course, a student:

- is able to prepare data for a given non-linear model
- train en test a non-linear model
- evaluate the quality of a trained model

Structure of the course

Table 1: Course overview

Week nr.	Module name	Readings
2	Onboarding and Introduction to the Course	Provost / Fawcett Ch.3
3-4	Lazy Learning with kNN	Provost / Fawcett Ch.6 + 7
5-6	Probabilistic Learning with Naive Bayes Classification	Provost / Fawcett Ch.9
7	Project Application	

Through the whole of the program you'll be cooperating within a team where you will combine and compare the results of the different case studies. At the end of the course you will present with your team what you have learned from analyzing and comparing the different case studies.

About the author



Witek ten Hove is a senior instructor and researcher at [HAN University of Applied Sciences](#). His main areas of expertise are Data en Web Technologies.

Through his extensive business experience in Finance and International Trade and thorough knowledge of modern data technologies, he is able to make connections between technology and business. As an open source evangelist he firmly believe in the power of knowledge sharing. His mission is to inspire business professionals and help them exploit the full potential of smart technologies.

He is the owner of [Ten Hove Business Data Solutions](#), a consultancy and training company helping organizations to achieve maximum business value through data driven solutions.

1 Setting up your data science environment

1.1 Working with Quarto

1.2 Working with Git and Github

1.3 Using Python virtual environments

2 Lazy learning with k-Nearest Neighbors

2.1 Business Case: Diagnosing Breast Cancer

Breast cancer is the top cancer in women both in the developed and the developing world. In the Netherlands it is the most pervasive form of cancer (“WHO | Cancer Country Profiles 2020” n.d.). In order to improve breast cancer outcome and survival early detection remains the most important instrument for breast cancer control. If machine learning could automate the identification of cancer, it would improve efficiency of the detection process and might also increase its effectiveness by providing greater detection accuracy.

2.2 Data Understanding

The data we will be using comes from the University of Wisconsin and is available online as an open source dataset (“UCI Machine Learning Repository: Breast Cancer Wisconsin (Diagnostic) Data Set” n.d.). It includes measurements from digitized images from from fine-needle aspirates of breast mass. The values represent cell nuclei features.

For convenience the data in csv format is stored on Github. We can access it directly using a function for reading csv from the `pandas` library

```
url = "https://raw.githubusercontent.com/businessdatasolutions/courses/main/data%20mining/"
rawDF = pd.read_csv(url)
```

Using the `info()` function we can have some basic information about the dataset.

```
rawDF.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 32 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    569 non-null   int64
1   diagnosis             569 non-null   object
```

2	radius_mean	569	non-null	float64
3	texture_mean	569	non-null	float64
4	perimeter_mean	569	non-null	float64
5	area_mean	569	non-null	float64
6	smoothness_mean	569	non-null	float64
7	compactness_mean	569	non-null	float64
8	concavity_mean	569	non-null	float64
9	points_mean	569	non-null	float64
10	symmetry_mean	569	non-null	float64
11	dimension_mean	569	non-null	float64
12	radius_se	569	non-null	float64
13	texture_se	569	non-null	float64
14	perimeter_se	569	non-null	float64
15	area_se	569	non-null	float64
16	smoothness_se	569	non-null	float64
17	compactness_se	569	non-null	float64
18	concavity_se	569	non-null	float64
19	points_se	569	non-null	float64
20	symmetry_se	569	non-null	float64
21	dimension_se	569	non-null	float64
22	radius_worst	569	non-null	float64
23	texture_worst	569	non-null	float64
24	perimeter_worst	569	non-null	float64
25	area_worst	569	non-null	float64
26	smoothness_worst	569	non-null	float64
27	compactness_worst	569	non-null	float64
28	concavity_worst	569	non-null	float64
29	points_worst	569	non-null	float64
30	symmetry_worst	569	non-null	float64
31	dimension_worst	569	non-null	float64

dtypes: float64(30), int64(1), object(1)
memory usage: 142.4+ KB

The dataset has `python rawDF.shape()[0]` variables (columns).

2.3 Preparation

The first variable, `id`, contains unique patient IDs. The IDs do not contain any relevant information for making predictions, so we will delete it from the dataset.


```
cleanDF = rawDF.drop(['id'], axis=1)
cleanDF.head()
```

	diagnosis	radius_mean	...	symmetry_worst	dimension_worst
0	B	12.32	...	0.2827	0.06771
1	B	10.60	...	0.2940	0.07587
2	B	11.04	...	0.2998	0.07881
3	B	11.28	...	0.2102	0.06784
4	B	15.19	...	0.2487	0.06766

[5 rows x 31 columns]

The variable named **diagnosis** contains the outcomes we would like to predict - 'B' for 'Benign' and 'M' for 'Malignant'. The variable we would like to predict is called the 'label'. We can look at the counts both outcomes, using the `value_counts()` function. When we set the `normalize` setting to `True` we get the the proportions.

```
cntDiag = cleanDF['diagnosis'].value_counts()
propDiag = cleanDF['diagnosis'].value_counts(normalize=True)
cntDiag
```

```
B    357
M    212
Name: diagnosis, dtype: int64
```

```
propDiag
```

```
B    0.627417
M    0.372583
Name: diagnosis, dtype: float64
```

Looking again at the results from the `info()` function we notice that The variable **diagnosis** is coded as text (`object`). Many models require that the label is of type `category`. The 'pandas library contains a function that can transform a `object` type to `category`.

```
catType = CategoricalDtype(categories=["B", "M"], ordered=False)
cleanDF['diagnosis'] = cleanDF['diagnosis'].astype(catType)
cleanDF['diagnosis']
```

```

0      B
1      B
2      B
3      B
4      B
..
564    B
565    B
566    M
567    B
568    M
Name: diagnosis, Length: 569, dtype: category
Categories (2, object): ['B', 'M']

```

The features consist of three different measurements of ten characteristics. We will take three characteristics and have a closer look.

```
cleanDF[['radius_mean', 'area_mean', 'smoothness_mean']].describe()
```

	radius_mean	area_mean	smoothness_mean
count	569.000000	569.000000	569.000000
mean	14.127292	654.889104	0.096360
std	3.524049	351.914129	0.014064
min	6.981000	143.500000	0.052630
25%	11.700000	420.300000	0.086370
50%	13.370000	551.100000	0.095870
75%	15.780000	782.700000	0.105300
max	28.110000	2501.000000	0.163400

You'll notice that the three variables have very different ranges and as a consequence **area_mean** will have a larger impact on the distance calculation than the **smoothness_mean**. This could potentially cause problems for modeling. To solve this we'll apply normalization to rescale all features to a standard range of values.

We will write our own normalization function.

```

def normalize(x):
    return((x - min(x)) / (max(x) - min(x))) # distance of item value - minimum vector value

testSet1 = np.arange(1,6)
testSet2 = np.arange(1,6) * 10

```

```
print(f'testSet1: {testSet1}\n')
```

```
testSet1: [1 2 3 4 5]
```

```
print(f'testSet2: {testSet1}\n')
```

```
testSet2: [1 2 3 4 5]
```

```
print(f'Normalized testSet1: {normalize(testSet1)}\n')
```

```
Normalized testSet1: [0.  0.25 0.5  0.75 1.  ]
```

```
print(f'Normalized testSet2: {normalize(testSet2)}\n')
```

```
Normalized testSet2: [0.  0.25 0.5  0.75 1.  ]
```

and apply it to all the numerical variables in the dataframe.

```
excluded = ['diagnosis'] # list of columns to exclude
X = cleanDF.loc[:, ~cleanDF.columns.isin(excluded)]
X = X.apply(normalize, axis=0)
X[['radius_mean', 'area_mean', 'smoothness_mean']].describe()
```

	radius_mean	area_mean	smoothness_mean
count	569.000000	569.000000	569.000000
mean	0.338222	0.216920	0.394785
std	0.166787	0.149274	0.126967
min	0.000000	0.000000	0.000000
25%	0.223342	0.117413	0.304595
50%	0.302381	0.172895	0.390358
75%	0.416442	0.271135	0.475490
max	1.000000	1.000000	1.000000

When we take the variables we selected earlier and look at the summary parameters again, we'll see that the normalization was successful.

We can now split our data into training and test sets.

```
y = cleanDF['diagnosis']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=123,
```

Now we can train and evaluate our kNN model.

2.4 Modeling and Evaluation

To train the knn model we only need one single function from the `sklearn` library. It takes the set with training features and the set with training label and fits a model to the training data. The trained model is applied to the set with test features and the `predict()` function gives back a set of predictions.

```
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)
# make predictions on the test set
```

`KNeighborsClassifier()`

```
y_pred = knn.predict(X_test)
```

Now that we have a set of predicted labels we can compare these with the actual labels. A confusion table shows how well the model performed.

Here is our own table:

```
cm = confusion_matrix(y_test, y_pred, labels=knn.classes_)
cm
```

```
array([[106,  1],
       [ 2, 62]])
```


```
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=knn.classes_)
disp.plot()
```

		True class		Measures
		Positive	Negative	
Predicted class	Positive	True positive <i>TP</i>	False positive <i>FP</i>	Positive predictive value (PPV) $\frac{TP}{TP+FP}$
	Negative	False negative <i>FN</i>	True negative <i>TN</i>	Negative predictive value (NPV) $\frac{TN}{FN+TN}$
Measures		Sensitivity $\frac{TP}{TP+FN}$	Specificity $\frac{TN}{FP+TN}$	Accuracy $\frac{TP+TN}{TP+FP+FN+TN}$

Figure 2.1: Standard confusion table. Taken from: <https://emj.bmj.com/content/emjmed/36/7/431/F1.large.jpg>

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x1ae7dac20>

```
plt.show()
```



Questions:

1. *How would you assess the overall performance of the model?*
2. *What would you consider as more costly: high false negatives or high false positives levels? Why?*

3 Probabilistic Learning with Naive Bayes Classification

3.1 Business Case: Filtering Spam

In 2020 spam accounted for more than 50% of total e-mail traffic (“Spam Statistics: Spam e-Mail Traffic Share 2019” n.d.). This illustrates the value of a good spam filter. Naive Bayes spam filtering is a standard technique for handling spam. It is one of the oldest ways of doing spam filtering, with roots in the 1990s.

3.2 Data Understanding

The data you’ll be using comes from the SMS Spam Collection (“UCI Machine Learning Repository: SMS Spam Collection Data Set” n.d.). It contains a set of SMS messages that are labeled ‘ham’ or ‘spam’. and is a standard data set for testing spam filtering methods.

```
url = "datasets/smsspam.csv"
rawDF = pd.read_csv(url)
rawDF.head()
```

	type	text
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...

The variable `type` is of class `object` which in Python refers to text. As this variable indicates whether the message belongs to the category `ham` or `spam` it is better to convert it to a `category` variable.

```
catType = CategoricalDtype(categories=["ham", "spam"], ordered=False)
rawDF.type = rawDF.type.astype(catType)
```

```
rawDF.type
```

```
0      ham
1      ham
2     spam
3      ham
4      ham
...
5567   spam
5568   ham
5569   ham
5570   ham
5571   ham
Name: type, Length: 5572, dtype: category
Categories (2, object): ['ham', 'spam']
```

To see how the types of sms messages are distributed you can compare the counts for each category.

```
rawDF.type.value_counts()
```

```
ham      4825
spam      747
Name: type, dtype: int64
```

Often you'll prefer the relative counts.

```
rawDF.type.value_counts(normalize=True)
```

```
ham      0.865937
spam      0.134063
Name: type, dtype: float64
```

You can also visually inspect the data by creating wordclouds for each sms type.

```
# Generate a word cloud image]
hamText = ' '.join([Text for Text in rawDF[rawDF['type']=='ham']['text']])
spamText = ' '.join([Text for Text in rawDF[rawDF['type']=='spam']['text']])
colorListHam=['#e9f6fb','#92d2ed','#2195c5']
```


3.3 Preparation

After you’ve glimpsed over the data and have a certain understanding of its structure and content, you are now ready to prepare the data for further processing. For the naive bayes model you’ll need to have a dataframe with wordcounts. To save on computation time you can set a limit on the number of features (columns) in the wordsDF dataframe.

```
vectorizer = TfidfVectorizer(max_features=1000)
vectors = vectorizer.fit_transform(rawDF.text)
wordsDF = pd.DataFrame(vectors.toarray(), columns=vectorizer.get_feature_names_out())
wordsDF.head()
```

	000	03	04	0800	08000839402	...	your	yours	yourself	yr	yup
0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0

[5 rows x 1000 columns]

The counts are normalized in such a way that the words that are most likely to have predictive power get heavier weights. For instance stopword like “a” and “for” most probably will equally likely feature in spam as in ham messages. Therefore these words will be assigned lower normalized counts.

Before we start modeling we need to split all datasets into *train* and *test* sets. The function `train_test_split()` can be used to create balanced splits of the data. In this case we’ll create a 75/25% split.

```
xTrain, xTest, yTrain, yTest = train_test_split(wordsDF, rawDF.type)
```

3.4 Modeling and Evaluation

We have now everything in place to start training our model and evaluate against our test dataset. The `MultinomialNB().fit()` function is part of the `scikit learn` package. It takes in the features and labels of our training dataset and returns a trained naive bayes model.

```

bayes = MultinomialNB()
bayes.fit(xTrain, yTrain)

```

MultinomialNB()

The model can be applied to the test features using the `predict()` function which generates a array of predictions. Using a confusion matrix we can analyze the performance of our model.

		True class		Measures
		Positive	Negative	
Predicted class	Positive	True positive <i>TP</i>	False positive <i>FP</i>	Positive predictive value (PPV) $\frac{TP}{TP+FP}$
	Negative	False negative <i>FN</i>	True negative <i>TN</i>	Negative predictive value (NPV) $\frac{TN}{FN+TN}$
Measures		Sensitivity $\frac{TP}{TP+FN}$	Specificity $\frac{TN}{FP+TN}$	Accuracy $\frac{TP+TN}{TP+FP+FN+TN}$

Figure 3.1: Standard diffusion table. Taken from: <https://emj.bmj.com/content/emered/36/7/431/F1.large.jp>

```

yPred = bayes.predict(xTest)
yTrue = yTest

accuracyScore = accuracy_score(yTrue, yPred)
print(f'Accuracy: {accuracyScore}')

```

Accuracy: 0.9748743718592965

```

matrix = confusion_matrix(yTrue, yPred)
labelNames = pd.Series(['ham', 'spam'])
pd.DataFrame(matrix,
              columns='Predicted ' + labelNames,
              index='Is ' + labelNames)

```

	Predicted ham	Predicted spam
Is ham	1220	0
Is spam	35	138

Questions:

1. What do you think is the role of the **alpha** parameter in the *MultinomialNB()* function?
2. How would you assess the overall performance of the model?
3. What would you consider as more costly: high false negatives or high false positives levels? Why?

References

- “Spam Statistics: Spam e-Mail Traffic Share 2019.” n.d. *Statista*. Accessed January 10, 2021. <https://www.statista.com/statistics/420391/spam-email-traffic-share/>.
- “UCI Machine Learning Repository: Breast Cancer Wisconsin (Diagnostic) Data Set.” n.d. Accessed January 7, 2021. [https://archive.ics.uci.edu/ml/datasets/breast+cancer+wisconsin+\(diagnostic\)](https://archive.ics.uci.edu/ml/datasets/breast+cancer+wisconsin+(diagnostic)).
- “UCI Machine Learning Repository: SMS Spam Collection Data Set.” n.d. Accessed January 9, 2021. <https://archive.ics.uci.edu/ml/datasets/sms+spam+collection>.
- “WHO | Cancer Country Profiles 2020.” n.d. *WHO*. Accessed January 7, 2021. <http://www.who.int/cancer/country-profiles/en/>.